

lecture_02.py



```
1 from execute_util import text, link, image
2 from facts import a100_flop_per_sec, h100_flop_per_sec
3 import torch.nn.functional as F
4 import timeit
5 import torch
6 from typing import Iterable
7 from torch import nn
8 import numpy as np
9 from lecture_util import article_link
10 from jaxtyping import Float
11 from einops import rearrange, einsum, reduce
12 from references import zero_2019
13
14
15 def main():
16     Last lecture: overview, tokenization
17
18     Overview of this lecture:
19     • We will discuss all the primitives needed to train a model.
20     • We will go bottom-up from tensors to models to optimizers to the training loop.
21     • We will pay close attention to efficiency (use of resources).
22
23     In particular, we will account for two types of resources:
24     • Memory (GB)
25     • Compute (FLOPs)
26
27     motivating_questions()
28
29     We will not go over the Transformer.
30     There are excellent expositions:
31     Assignment 1 handout
32     Mathematical description
33     Illustrated Transformer
34     Illustrated GPT-2
35     Instead, we'll work with simpler models.
36
37     What knowledge to take away:
38     • Mechanics: straightforward (just PyTorch)
39     • Mindset: resource accounting (remember to do it)
40     • Intuitions: broad strokes (no large models)
41
42     Memory accounting
43     tensors_basics()
44     tensors_memory()
45
46     Compute accounting
47     tensors_on_gpus()
48     tensor_operations()
49     tensor_einops()
50     tensor_operations_flops()
51     gradients_basics()
52     gradients_flops()
53
54     Models
55     module_parameters()
56     custom_model()
57
58     Training loop and best practices
```

```

59  note_about_randomness()
60  data_loading()
61
62  optimizer()
63  train_loop()
64  checkpointing()
65  mixed_precision_training()
66
67
68  def motivating_questions():
69      Let's do some napkin math.
70
71      Question: How long would it take to train a 70B parameter model on 15T tokens on 1024 H100s?
72      total_flops = 6 * 70e9 * 15e12 # @inspect total_flops
73      assert h100_flop_per_sec == 1979e12 / 2
74      mfu = 0.5
75      flops_per_day = h100_flop_per_sec * mfu * 1024 * 60 * 60 * 24 # @inspect flops_per_day
76      days = total_flops / flops_per_day # @inspect days
77
78      Question: What's the largest model that can you can train on 8 H100s using AdamW (naively)?
79      h100_bytes = 80e9 # @inspect h100_bytes
80      bytes_per_parameter = 4 + 4 + (4 + 4) # parameters, gradients, optimizer state @inspect bytes_per_parameter
81      num_parameters = (h100_bytes * 8) / bytes_per_parameter # @inspect num_parameters
82      Caveat 1: we are naively using float32 for parameters and gradients. We could also use bf16 for parameters
      and gradients (2 + 2) and keep an extra float32 copy of the parameters (4). This doesn't save memory, but is
      faster. [Rajbhandari+ 2019]
83      Caveat 2: activations are not accounted for (depends on batch size and sequence length).
84
85      This is a rough back-of-the-envelope calculation.
86
87
88  def tensors_basics():
89      Tensors are the basic building block for storing everything: parameters, gradients, optimizer state, data,
      activations.
90      [PyTorch docs on tensors]
91
92      You can create tensors in multiple ways:
93      x = torch.tensor([[1., 2, 3], [4, 5, 6]]) # @inspect x
94      x = torch.zeros(4, 8) # 4x8 matrix of all zeros @inspect x
95      x = torch.ones(4, 8) # 4x8 matrix of all ones @inspect x
96      x = torch.randn(4, 8) # 4x8 matrix of iid Normal(0, 1) samples @inspect x
97
98      Allocate but don't initialize the values:
99      x = torch.empty(4, 8) # 4x8 matrix of uninitialized values @inspect x
100      ...because you want to use some custom logic to set the values later
101      nn.init.trunc_normal_(x, mean=0, std=1, a=-2, b=2) # @inspect x
102
103
104  def tensors_memory():
105      Almost everything (parameters, gradients, activations, optimizer states) are stored as floating point numbers.
106
107
108  float32
109  [Wikipedia]
110  IEEE 754 single-precision 32-bit float
111
112  sign | exponent (8 bit) | fraction (23 bit)
113  0 0 1 1 1 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
114  31 30 23 22 0

```

The float32 data type (also known as fp32 or single precision) is the default. Traditionally, in scientific computing, float32 is the baseline; you could use double precision (float64) in some cases.

In deep learning, you can be a lot sloppier.

Let's examine memory usage of these tensors.

Memory is determined by the (i) number of values and (ii) data type of each value.

```


116 x = torch.zeros(4, 8) # @inspect x
117 assert x.dtype == torch.float32 # Default type
118 assert x.numel() == 4 * 8
119 assert x.element_size() == 4 # Float is 4 bytes
120 assert get_memory_usage(x) == 4 * 8 * 4 # 128 bytes
121
122 One matrix in the feedforward layer of GPT-3:
123 assert get_memory_usage(torch.empty(12288 * 4, 12288)) == 2304 * 1024 * 1024 # 2.3 GB
124 ...which is a lot!

```

float16

[Wikipedia]

IEEE half-precision 16-bit float



```


129 The float16 data type (also known as fp16 or half precision) cuts down the memory.
130 x = torch.zeros(4, 8, dtype=torch.float16) # @inspect x
131 assert x.element_size() == 2
132 However, the dynamic range (especially for small numbers) isn't great.
133 x = torch.tensor([1e-8], dtype=torch.float16) # @inspect x
134 assert x == 0 # Underflow!
135 If this happens when you train, you can get instability.

```

bfloat16

[Wikipedia]

bfloat16



```

140 Google Brain developed bfloat (brain floating point) in 2018 to address this issue.
141 bfloat16 uses the same memory as float16 but has the same dynamic range as float32!
142 The only catch is that the resolution is worse, but this matters less for deep learning.
143 x = torch.tensor([1e-8], dtype=torch.bfloat16) # @inspect x
144 assert x != 0 # No underflow!

```

Let's compare the dynamic ranges and memory usage of the different data types:

```

147 float32_info = torch.finfo(torch.float32) # @inspect float32_info
148 float16_info = torch.finfo(torch.float16) # @inspect float16_info
149 bfloat16_info = torch.finfo(torch.bfloat16) # @inspect bfloat16_info

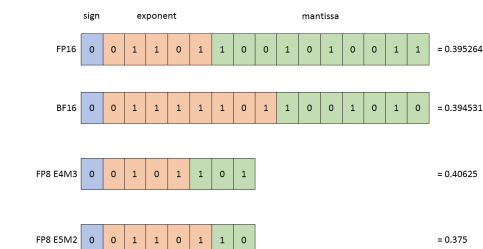
```

fp8

```

152 In 2022, FP8 was standardized, motivated by machine learning workloads.
153 https://docs.nvidia.com/deeplearning/transformer-engine/user-guide/examples/fp8_primer.html

```



```

155 H100s support two variants of FP8: E4M3 (range [-448, 448]) and E5M2 ([-57344, 57344]).
156 Reference: [Micikevicius+ 2022]

```

Implications on training:

- Training with float32 works, but requires lots of memory.
- Training with fp8, float16 and even bfloat16 is risky, and you can get instability.
- Solution (later): use mixed precision training, see [mixed_precision_training](#)

163

164 `def tensors_on_gpus():`


165 By default, tensors are stored in CPU memory.

166 `x = torch.zeros(32, 32)`167 `assert x.device == torch.device("cpu")`

168

169 However, in order to take advantage of the massive parallelism of GPUs, we need to move them to GPU memory.

170



171

172 Let's first see if we have any GPUs.

173 `if not torch.cuda.is_available():`174 `return`

175

176 `num_gpus = torch.cuda.device_count() # @inspect num_gpus`177 `for i in range(num_gpus):`178 `properties = torch.cuda.get_device_properties(i) # @inspect properties`

179

180 `memory_allocated = torch.cuda.memory_allocated() # @inspect memory_allocated`

181

182 `text("Move the tensor to GPU memory (device 0).")`183 `y = x.to("cuda:0")`184 `assert y.device == torch.device("cuda", 0)`

185

186 `text("Or create a tensor directly on the GPU:")`187 `z = torch.zeros(32, 32, device="cuda:0")`

188

189 `new_memory_allocated = torch.cuda.memory_allocated() # @inspect new_memory_allocated`190 `memory_used = new_memory_allocated - memory_allocated # @inspect memory_used`191 `assert memory_used == 2 * (32 * 32 * 4) # 2 32x32 matrices of 4-byte floats`

192

193

194

195 `def tensor_operations():`

196 Most tensors are created from performing operations on other tensors.

197 Each operation has some memory and compute consequence.

198

199 `tensor_storage()`200 `tensor_slicing()`201 `tensor_elementwise()`202 `tensor_matmul()`

203

204

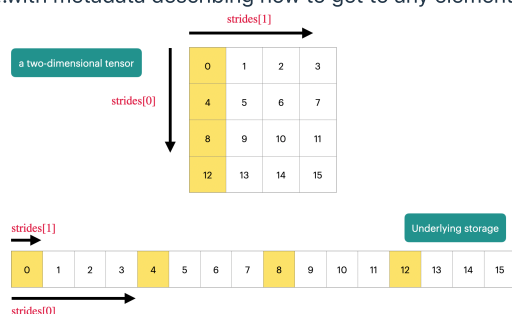
205 `def tensor_storage():`

206 What are tensors in PyTorch?

207 PyTorch tensors are pointers into allocated memory

208 ...with metadata describing how to get to any element of the tensor.

209



210

210 [\[PyTorch docs\]](#)211 `x = torch.tensor([`212 `[0., 1, 2, 3],`

```


213         [4, 5, 6, 7],
214         [8, 9, 10, 11],
215         [12, 13, 14, 15],
216     ])
217
218     To go to the next row (dim 0), skip 4 elements in storage.
219     assert x.stride(0) == 4
220
221     To go to the next column (dim 1), skip 1 element in storage.
222     assert x.stride(1) == 1
223
224     To find an element:
225     r, c = 1, 2
226     index = r * x.stride(0) + c * x.stride(1) # @inspect index
227     assert index == 6
228
229
230 def tensor_slicing():
231     x = torch.tensor([[1., 2, 3], [4, 5, 6]]) # @inspect x
232
233     Many operations simply provide a different view of the tensor.
234     This does not make a copy, and therefore mutations in one tensor affects the other.
235
236     Get row 0:
237     y = x[0] # @inspect y
238     assert torch.equal(y, torch.tensor([1., 2, 3]))
239     assert same_storage(x, y)
240
241     Get column 1:
242     y = x[:, 1] # @inspect y
243     assert torch.equal(y, torch.tensor([2, 5]))
244     assert same_storage(x, y)
245
246     View 2x3 matrix as 3x2 matrix:
247     y = x.view(3, 2) # @inspect y
248     assert torch.equal(y, torch.tensor([[1, 2], [3, 4], [5, 6]]))
249     assert same_storage(x, y)
250
251     Transpose the matrix:
252     y = x.transpose(1, 0) # @inspect y
253     assert torch.equal(y, torch.tensor([[1, 4], [2, 5], [3, 6]]))
254     assert same_storage(x, y)
255
256     Check that mutating x also mutates y.
257     x[0][0] = 100 # @inspect x, @inspect y
258     assert y[0][0] == 100
259
260     Note that some views are non-contiguous entries, which means that further views aren't possible.
261     x = torch.tensor([[1., 2, 3], [4, 5, 6]]) # @inspect x
262     y = x.transpose(1, 0) # @inspect y
263     assert not y.is_contiguous()
264     try:
265         y.view(2, 3)
266         assert False
267     except RuntimeError as e:
268         assert "view size is not compatible with input tensor's size and stride" in str(e)
269
270     One can enforce a tensor to be contiguous first:
271     y = x.transpose(1, 0).contiguous().view(2, 3) # @inspect y
272     assert not same_storage(x, y)
273     Views are free, copying take both (additional) memory and compute.
274
275
276 def tensor_elementwise():

```

```

277 These operations apply some operation to each element of the tensor
278 ...and return a (new) tensor of the same shape.
279
280 x = torch.tensor([1, 4, 9])
281 assert torch.equal(x.pow(2), torch.tensor([1, 16, 81]))
282 assert torch.equal(x.sqrt(), torch.tensor([1, 2, 3]))
283 assert torch.equal(x.rsqrt(), torch.tensor([1, 1 / 2, 1 / 3])) # i -> 1/sqrt(x_i)
284
285 assert torch.equal(x + x, torch.tensor([2, 8, 18]))
286 assert torch.equal(x * 2, torch.tensor([2, 8, 18]))
287 assert torch.equal(x / 0.5, torch.tensor([2, 8, 18]))
288
289 triu takes the upper triangular part of a matrix.
290 x = torch.ones(3, 3).triu() # @inspect x
291 assert torch.equal(x, torch.tensor([
292     [1, 1, 1],
293     [0, 1, 1],
294     [0, 0, 1]],
295 ))
296 This is useful for computing an causal attention mask, where  $M[i, j]$  is the contribution of  $i$  to  $j$ .
297
298
299 def tensor_matmul():
300     Finally, the bread and butter of deep learning: matrix multiplication.
301     x = torch.ones(16, 32)
302     w = torch.ones(32, 2)
303     y = x @ w
304     assert y.size() == torch.Size([16, 2])
305
306 In general, we perform operations for every example in a batch and token in a sequence.
307

```



```

308 x = torch.ones(4, 8, 16, 32)
309 w = torch.ones(32, 2)
310 y = x @ w
311 assert y.size() == torch.Size([4, 8, 16, 2])
312 In this case, we iterate over values of the first 2 dimensions of x and multiply by w.
313
314
315 def tensor_einops():
316     einops_motivation()
317
318 Einops is a library for manipulating tensors where dimensions are named.
319 It is inspired by Einstein summation notation (Einstein, 1916).
320 \[Einops tutorial\]
321
322 jaxtyping_basics()
323 einops_einsum()
324 einops_reduce()
325 einops_rearrange()
326
327
328 def einops_motivation():
329     Traditional PyTorch code:
330     x = torch.ones(2, 2, 3) # batch, sequence, hidden @inspect x
331     y = torch.ones(2, 2, 3) # batch, sequence, hidden @inspect y
332     z = x @ y.transpose(-2, -1) # batch, sequence, sequence @inspect z

```

```

333     Easy to mess up the dimensions (what is -2, -1?)...
334
335
336 def jaxtyping_basics():
337     How do you keep track of tensor dimensions?
338
339     Old way:
340     x = torch.ones(2, 2, 1, 3) # batch seq heads hidden @inspect x
341
342     New (jaxtyping) way:
343     x: Float[torch.Tensor, "batch seq heads hidden"] = torch.ones(2, 2, 1, 3) # @inspect x
344     Note: this is just documentation (no enforcement).
345
346
347 def einops_einsum():
348     Einsum is generalized matrix multiplication with good bookkeeping.
349
350     Define two tensors:
351     x: Float[torch.Tensor, "batch seq1 hidden"] = torch.ones(2, 3, 4) # @inspect x
352     y: Float[torch.Tensor, "batch seq2 hidden"] = torch.ones(2, 3, 4) # @inspect y
353
354     Old way:
355     z = x @ y.transpose(-2, -1) # batch, sequence, sequence @inspect z
356
357     New (einops) way:
358     z = einsum(x, y, "batch seq1 hidden, batch seq2 hidden -> batch seq1 seq2") # @inspect z
359     Dimensions that are not named in the output are summed over.
360
361     Or can use ... to represent broadcasting over any number of dimensions:
362     z = einsum(x, y, "... seq1 hidden, ... seq2 hidden -> ... seq1 seq2") # @inspect z
363
364
365 def einops_reduce():
366     You can reduce a single tensor via some operation (e.g., sum, mean, max, min).
367     x: Float[torch.Tensor, "batch seq hidden"] = torch.ones(2, 3, 4) # @inspect x
368
369     Old way:
370     y = x.mean(dim=-1) # @inspect y
371
372     New (einops) way:
373     y = reduce(x, "... hidden -> ...", "sum") # @inspect y
374
375
376 def einops_rearrange():
377     Sometimes, a dimension represents two dimensions
378     ...and you want to operate on one of them.
379
380     x: Float[torch.Tensor, "batch seq total_hidden"] = torch.ones(2, 3, 8) # @inspect x
381     ...where total_hidden is a flattened representation of heads * hidden1
382     w: Float[torch.Tensor, "hidden1 hidden2"] = torch.ones(4, 4)
383
384     Break up total_hidden into two dimensions (heads and hidden1):
385     x = rearrange(x, "... (heads hidden1) -> ... heads hidden1", heads=2) # @inspect x
386
387     Perform the transformation by w:
388     x = einsum(x, w, "... hidden1, hidden1 hidden2 -> ... hidden2") # @inspect x
389
390     Combine heads and hidden2 back together:
391     x = rearrange(x, "... heads hidden2 -> ... (heads hidden2)") # @inspect x
392
393
394 def tensor_operations_flops():
395     Having gone through all the operations, let us examine their computational cost.
396

```

A floating-point operation (FLOP) is a basic operation like addition ($x + y$) or multiplication ($x y$).

Two terribly confusing acronyms (pronounced the same!):

- FLOPs: floating-point operations (measure of computation done)
- FLOP/s: floating-point operations per second (also written as FLOPS), which is used to measure the speed of hardware.

Intuitions

Training GPT-3 (2020) took $3.14e23$ FLOPs. [\[article\]](#)

Training GPT-4 (2023) is speculated to take $2e25$ FLOPs [\[article\]](#)

US executive order: any foundation model trained with $\geq 1e26$ FLOPs must be reported to the government (revoked in 2025)

A100 has a peak performance of 312 teraFLOP/s [\[spec\]](#)

```
assert a100_flop_per_sec == 312e12
```

H100 has a peak performance of 1979 teraFLOP/s with sparsity, 50% without [\[spec\]](#)

```
assert h100_flop_per_sec == 1979e12 / 2
```

8 H100s for 2 weeks:

```
total_flops = 8 * (60 * 60 * 24 * 7) * h100_flop_per_sec # @inspect total_flops
```

Linear model

As motivation, suppose you have a linear model.

- We have n points
- Each point is d -dimensional
- The linear model maps each d -dimensional vector to a k outputs

```
if torch.cuda.is_available():
    B = 16384 # Number of points
    D = 32768 # Dimension
    K = 8192 # Number of outputs
else:
    B = 1024
    D = 256
    K = 64
```

```
device = get_device()
x = torch.ones(B, D, device=device)
w = torch.randn(D, K, device=device)
y = x @ w
```

We have one multiplication ($x[i][j] * w[j][k]$) and one addition per (i, j, k) triple.

```
actual_num_flops = 2 * B * D * K # @inspect actual_num_flops
```

FLOPs of other operations

- Elementwise operation on a $m \times n$ matrix requires $O(mn)$ FLOPs.
- Addition of two $m \times n$ matrices requires $m \times n$ FLOPs.

In general, no other operation that you'd encounter in deep learning is as expensive as matrix multiplication for large enough matrices.

Interpretation:

- B is the number of data points
 - $(D K)$ is the number of parameters
 - FLOPs for forward pass is $2 \times (\text{\# tokens}) \times (\text{\# parameters})$
- It turns out this generalizes to Transformers (to a first-order approximation).

How do our FLOPs calculations translate to wall-clock time (seconds)?

Let us time it!

```
actual_time = time_matmul(x, w) # @inspect actual_time
actual_flop_per_sec = actual_num_flops / actual_time # @inspect actual_flop_per_sec
```

Each GPU has a specification sheet that reports the peak performance.

- A100 [\[spec\]](#)

- H100 [spec]

Note that the FLOP/s depends heavily on the data type!

```
promised_flop_per_sec = get_promised_flop_per_sec(device, x.dtype) # @inspect promised_flop_per_sec
```

Model FLOPs utilization (MFU)

Definition: (actual FLOP/s) / (promised FLOP/s) [ignore communication/overhead]

```
mfu = actual_flop_per_sec / promised_flop_per_sec # @inspect mfu
```

Usually, MFU of ≥ 0.5 is quite good (and will be higher if matmuls dominate)

Let's do it with bfloat16:

```
x = x.to(torch.bfloat16)
```

```
w = w.to(torch.bfloat16)
```

```
bf16_actual_time = time_matmul(x, w) # @inspect bf16_actual_time
```

```
bf16_actual_flop_per_sec = actual_num_flops / bf16_actual_time # @inspect bf16_actual_flop_per_sec
```

```
bf16_promised_flop_per_sec = get_promised_flop_per_sec(device, x.dtype) # @inspect bf16_promised_flop_per_sec
```

```
bf16_mfu = bf16_actual_flop_per_sec / bf16_promised_flop_per_sec # @inspect bf16_mfu
```

Note: comparing bfloat16 to float32, the actual FLOP/s is higher.

The MFU here is rather low, probably because the promised FLOPs is a bit optimistic.

Summary

- Matrix multiplications dominate: (2 m n p) FLOPs
- FLOP/s depends on hardware (H100 >> A100) and data type (bfloat16 >> float32)
- Model FLOPs utilization (MFU): (actual FLOP/s) / (promised FLOP/s)

```
def gradients_basics():
```

So far, we've constructed tensors (which correspond to either parameters or data) and passed them through operations (forward).

Now, we're going to compute the gradient (backward).

As a simple example, let's consider the simple linear model:

$$y = 0.5 (x * w - 5)^2$$

Forward pass: compute loss

```
x = torch.tensor([1., 2, 3])
```

```
w = torch.tensor([1., 1, 1], requires_grad=True) # Want gradient
```

```
pred_y = x @ w
```

```
loss = 0.5 * (pred_y - 5).pow(2)
```

Backward pass: compute gradients

```
loss.backward()
```

```
assert loss.grad is None
```

```
assert pred_y.grad is None
```

```
assert x.grad is None
```

```
assert torch.equal(w.grad, torch.tensor([1, 2, 3]))
```

```
def gradients_flops():
```

Let us do count the FLOPs for computing gradients.

Revisit our linear model

```
if torch.cuda.is_available():
```

```
    B = 16384 # Number of points
```

```
    D = 32768 # Dimension
```

```
    K = 8192 # Number of outputs
```

```
else:
```

```
    B = 1024
```

```
    D = 256
```

```
    K = 64
```


```
device = get_device()
```

```
x = torch.ones(B, D, device=device)
```

```

519 w1 = torch.randn(D, D, device=device, requires_grad=True)
520 w2 = torch.randn(D, K, device=device, requires_grad=True)
521
522 Model: x --w1--> h1 --w2--> h2 -> loss
523 h1 = x @ w1
524 h2 = h1 @ w2
525 loss = h2.pow(2).mean()
526
527 Recall the number of forward FLOPs: tensor\_operations\_flops
528 • Multiply  $x[i][j] * w1[j][k]$ 
529 • Add to  $h1[i][k]$ 
530 • Multiply  $h1[i][j] * w2[j][k]$ 
531 • Add to  $h2[i][k]$ 
532 num_forward_flops = (2 * B * D * D) + (2 * B * D * K) # @inspect num_forward_flops
533
534 How many FLOPs is running the backward pass?
535 h1.retain_grad() # For debugging
536 h2.retain_grad() # For debugging
537 loss.backward()
538
539 Recall model: x --w1--> h1 --w2--> h2 -> loss
540
541 •  $h1.grad = d \text{ loss} / d h1$ 
542 •  $h2.grad = d \text{ loss} / d h2$ 
543 •  $w1.grad = d \text{ loss} / d w1$ 
544 •  $w2.grad = d \text{ loss} / d w2$ 
545
546 Focus on the parameter w2.
547 Invoke the chain rule.
548
549 num_backward_flops = 0 # @inspect num_backward_flops
550
551 w2.grad[j,k] = sum_i h1[i,j] * h2.grad[i,k]
552 assert w2.grad.size() == torch.Size([D, K])
553 assert h1.size() == torch.Size([B, D])
554 assert h2.grad.size() == torch.Size([B, K])
555 For each (i, j, k), multiply and add.
556 num_backward_flops += 2 * B * D * K # @inspect num_backward_flops
557
558 h1.grad[i,j] = sum_k w2[i,j] * h2.grad[i,k]
559 assert h1.grad.size() == torch.Size([B, D])
560 assert w2.size() == torch.Size([D, K])
561 assert h2.grad.size() == torch.Size([B, K])
562 For each (i, j, k), multiply and add.
563 num_backward_flops += 2 * B * D * K # @inspect num_backward_flops
564
565 This was for just w2 (D*K parameters).
566 Can do it for w1 (D*D parameters) as well (though don't need x.grad).
567 num_backward_flops += (2 + 2) * B * D * D # @inspect num_backward_flops
568
569 A nice graphical visualization: \[article\]
570

```



571

```

572 Putting it together:
573 • Forward pass: 2 (# data points) (# parameters) FLOPs
574 • Backward pass: 4 (# data points) (# parameters) FLOPs
575 • Total: 6 (# data points) (# parameters) FLOPs

```

```

576
577
578 def module_parameters():
579     input_dim = 16384
580     output_dim = 32
581
582     Model parameters are stored in PyTorch as nn.Parameter objects.
583     w = nn.Parameter(torch.randn(input_dim, output_dim))
584     assert isinstance(w, torch.Tensor) # Behaves like a tensor
585     assert type(w.data) == torch.Tensor # Access the underlying tensor
586

```

Parameter initialization

```

588
589 Let's see what happens.
590 x = nn.Parameter(torch.randn(input_dim))
591 output = x @ w # @inspect output
592 assert output.size() == torch.Size([output_dim])
593 Note that each element of output scales as sqrt(input_dim): -84.09819030761719.
594 Large values can cause gradients to blow up and cause training to be unstable.

```

```

595
596 We want an initialization that is invariant to input_dim.
597 To do that, we simply rescale by 1/sqrt(input_dim)
598 w = nn.Parameter(torch.randn(input_dim, output_dim) / np.sqrt(input_dim))
599 output = x @ w # @inspect output
600 Now each element of output is constant: -0.12920215725898743.

```

```

601
602 Up to a constant, this is Xavier initialization. \[paper\]\[stackexchange\]
603

```

```

604 To be extra safe, we truncate the normal distribution to [-3, 3] to avoid any chance of outliers.
605 w = nn.Parameter(nn.init.trunc_normal_(torch.empty(input_dim, output_dim), std=1 / np.sqrt(input_dim), a=-3, b=3))
606

```

```

607
608 def custom_model():
609     Let's build up a simple deep linear model using nn.Parameter.

```

```

610
611     D = 64 # Dimension
612     num_layers = 2
613     model = Cruncher(dim=D, num_layers=num_layers)
614
615     param_sizes = [
616         (name, param.numel())
617         for name, param in model.state_dict().items()
618     ]
619     assert param_sizes == [
620         ("layers.0.weight", D * D),
621         ("layers.1.weight", D * D),
622         ("final.weight", D),
623     ]
624     num_parameters = get_num_parameters(model)
625     assert num_parameters == (D * D) + (D * D) + D
626

```

```

627 Remember to move the model to the GPU.
628 device = get_device()
629 model = model.to(device)
630

```

```

631 Run the model on some data.
632 B = 8 # Batch size
633 x = torch.randn(B, D, device=device)
634 y = model(x)
635 assert y.size() == torch.Size([B])

```

```

636
637
638 class Linear(nn.Module):
639     """Simple linear layer."""
640     def __init__(self, input_dim: int, output_dim: int):
641         super().__init__()
642         self.weight = nn.Parameter(torch.randn(input_dim, output_dim) / np.sqrt(input_dim))
643
644     def forward(self, x: torch.Tensor) -> torch.Tensor:
645         return x @ self.weight
646
647
648 class Cruncher(nn.Module):
649     def __init__(self, dim: int, num_layers: int):
650         super().__init__()
651         self.layers = nn.ModuleList([
652             Linear(dim, dim)
653             for i in range(num_layers)
654         ])
655         self.final = Linear(dim, 1)
656
657     def forward(self, x: torch.Tensor) -> torch.Tensor:
658         # Apply linear layers
659         B, D = x.size()
660         for layer in self.layers:
661             x = layer(x)
662
663         # Apply final head
664         x = self.final(x)
665         assert x.size() == torch.Size([B, 1])
666
667         # Remove the last dimension
668         x = x.squeeze(-1)
669         assert x.size() == torch.Size([B])
670
671         return x
672
673
674 def get_batch(data: np.array, batch_size: int, sequence_length: int, device: str) -> torch.Tensor:
675     Sample batch_size random positions into data.
676     start_indices = torch.randint(len(data) - sequence_length, (batch_size,))
677     assert start_indices.size() == torch.Size([batch_size])
678
679     Index into the data.
680     x = torch.tensor([data[start:start + sequence_length] for start in start_indices])
681     assert x.size() == torch.Size([batch_size, sequence_length])
682

```

Pinned memory

By default, CPU tensors are in paged memory. We can explicitly pin.

```

686 if torch.cuda.is_available():
687     x = x.pin_memory()

```

This allows us to copy x from CPU into GPU asynchronously.

```

690 x = x.to(device, non_blocking=True)

```

This allows us to do two things in parallel (not done here):

- Fetch the next batch of data into CPU
- Process x on the GPU.

[\[article\]](#)

[\[article\]](#)

```

699 return x

```

```

700
701
702 def note_about_randomness():
703     Randomness shows up in many places: parameter initialization, dropout, data ordering, etc.
704     For reproducibility, we recommend you always pass in a different random seed for each use of randomness.
705     Determinism is particularly useful when debugging, so you can hunt down the bug.
706
707     There are three places to set the random seed which you should do all at once just to be safe.
708
709     # Torch
710     seed = 0
711     torch.manual_seed(seed)
712
713     # NumPy
714     import numpy as np
715     np.random.seed(seed)
716
717     # Python
718     import random
719     random.seed(seed)
720
721
722 def data_loading():
723     In language modeling, data is a sequence of integers (output by the tokenizer).
724
725     It is convenient to serialize them as numpy arrays (done by the tokenizer).
726     orig_data = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10], dtype=np.int32)
727     orig_data.tofile("data.npy")
728
729     You can load them back as numpy arrays.
730     Don't want to load the entire data into memory at once (LLaMA data is 2.8TB).
731     Use memmap to lazily load only the accessed parts into memory.
732     data = np.memmap("data.npy", dtype=np.int32)
733     assert np.array_equal(data, orig_data)
734
735     A data loader generates a batch of sequences for training.
736     B = 2 # Batch size
737     L = 4 # Length of sequence
738     x = get_batch(data, batch_size=B, sequence_length=L, device=get_device())
739     assert x.size() == torch.Size([B, L])
740
741
742 class SGD(torch.optim.Optimizer):
743     def __init__(self, params: Iterable[nn.Parameter], lr: float = 0.01):
744         super(SGD, self).__init__(params, dict(lr=lr))
745
746     def step(self):
747         for group in self.param_groups:
748             lr = group["lr"]
749             for p in group["params"]:
750                 grad = p.grad.data
751                 p.data -= lr * grad
752
753
754 class AdaGrad(torch.optim.Optimizer):
755     def __init__(self, params: Iterable[nn.Parameter], lr: float = 0.01):
756         super(AdaGrad, self).__init__(params, dict(lr=lr))
757
758     def step(self):
759         for group in self.param_groups:
760             lr = group["lr"]
761             for p in group["params"]:
762                 # Optimizer state
763                 state = self.state[p]

```

```

764         grad = p.grad.data
765
766         # Get squared gradients  $g2 = \sum_{i<t} g_i^2$ 
767         g2 = state.get("g2", torch.zeros_like(grad))
768
769         # Update optimizer state
770         g2 += torch.square(grad)
771         state["g2"] = g2
772
773         # Update parameters
774         p.data -= lr * grad / torch.sqrt(g2 + 1e-5)
775
776
777 def optimizer():
778     Recall our deep linear model.
779     B = 2
780     D = 4
781     num_layers = 2
782     model = Cruncher(dim=D, num_layers=num_layers).to(get_device())
783
784     Let's define the AdaGrad optimizer
785     • momentum = SGD + exponential averaging of grad
786     • AdaGrad = SGD + averaging by  $\text{grad}^2$ 
787     • RMSProp = AdaGrad + exponentially averaging of  $\text{grad}^2$ 
788     • Adam = RMSProp + momentum
789
790     AdaGrad: https://www.jmlr.org/papers/volume12/duchi11a/duchi11a.pdf
791     optimizer = AdaGrad(model.parameters(), lr=0.01)
792     state = model.state_dict() # @inspect state
793
794     Compute gradients
795     x = torch.randn(B, D, device=get_device())
796     y = torch.tensor([4., 5.], device=get_device())
797     pred_y = model(x)
798     loss = F.mse_loss(input=pred_y, target=y)
799     loss.backward()
800
801     Take a step
802     optimizer.step()
803     state = model.state_dict() # @inspect state
804
805     Free up the memory (optional)
806     optimizer.zero_grad(set_to_none=True)
807

```

Memory

```

809
810 # Parameters
811 num_parameters = (D * D * num_layers) + D # @inspect num_parameters
812 assert num_parameters == get_num_parameters(model)
813
814 # Activations
815 num_activations = B * D * num_layers # @inspect num_activations
816
817 # Gradients
818 num_gradients = num_parameters # @inspect num_gradients
819
820 # Optimizer states
821 num_optimizer_states = num_parameters # @inspect num_optimizer_states
822
823 # Putting it all together, assuming float32
824 total_memory = 4 * (num_parameters + num_activations + num_gradients + num_optimizer_states) # @inspect total_memory
825

```

Compute (for one step)

```
827 flops = 6 * B * num_parameters # @inspect flops
```

```
828
```

```
829
```

Transformers

```
830
```

831 The accounting for a Transformer is more complicated, but the same idea.

832 Assignment 1 will ask you to do that.

```
833
```

834 Blog post describing memory usage for Transformer training [\[article\]](#)

835 Blog post describing FLOPs for a Transformer: [\[article\]](#)

```
836
```

```
837
```

```
838 def train_loop():
```

```
839     Generate data from linear function with weights (0, 1, 2, ..., D-1).
```

```
840     D = 16
```

```
841     true_w = torch.arange(D, dtype=torch.float32, device=get_device())
```

```
842     def get_batch(B: int) -> tuple[torch.Tensor, torch.Tensor]:
```

```
843         x = torch.randn(B, D).to(get_device())
```

```
844         true_y = x @ true_w
```

```
845         return (x, true_y)
```

```
846
```

847 Let's do a basic run

```
848 train("simple", get_batch, D=D, num_layers=0, B=4, num_train_steps=10, lr=0.01)
```

```
849
```

850 Do some hyperparameter tuning

```
851 train("simple", get_batch, D=D, num_layers=0, B=4, num_train_steps=10, lr=0.1)
```

```
852
```

```
853
```

```
854 def train(name: str, get_batch,
```

```
855         D: int, num_layers: int,
```

```
856         B: int, num_train_steps: int, lr: float):
```

```
857     model = Cruncher(dim=D, num_layers=0).to(get_device())
```

```
858     optimizer = SGD(model.parameters(), lr=0.01)
```

```
859
```

```
860     for t in range(num_train_steps):
```

```
861         # Get data
```

```
862         x, y = get_batch(B=B)
```

```
863
```

```
864         # Forward (compute loss)
```

```
865         pred_y = model(x)
```

```
866         loss = F.mse_loss(pred_y, y)
```

```
867
```

```
868         # Backward (compute gradients)
```

```
869         loss.backward()
```

```
870
```

```
871         # Update parameters
```

```
872         optimizer.step()
```

```
873         optimizer.zero_grad(set_to_none=True)
```

```
874
```

```
875
```

```
876 def checkpointing():
```

877 Training language models take a long time and certainly will crash.

878 You don't want to lose all your progress.

```
879
```

880 During training, it is useful to periodically save your model and optimizer state to disk.

```
881
```

```
882 model = Cruncher(dim=64, num_layers=3).to(get_device())
```

```
883 optimizer = AdaGrad(model.parameters(), lr=0.01)
```

```
884
```

885 Save the checkpoint:

```
886 checkpoint = {
```

```
887     "model": model.state_dict(),
```

```
888     "optimizer": optimizer.state_dict(),
```

```
889 }
```

```
890 torch.save(checkpoint, "model_checkpoint.pt")
```

```

891
892 Load the checkpoint:
893 loaded_checkpoint = torch.load("model_checkpoint.pt")
894
895
896 def mixed_precision_training():
897     Choice of data type (float32, bfloat16, fp8) have tradeoffs.
898     • Higher precision: more accurate/stable, more memory, more compute
899     • Lower precision: less accurate/stable, less memory, less compute
900
901     How can we get the best of both worlds?
902
903     Solution: use float32 by default, but use {bfloat16, fp8} when possible.
904
905     A concrete plan:
906     • Use {bfloat16, fp8} for the forward pass (activations).
907     • Use float32 for the rest (parameters, gradients).
908
909     • Mixed precision training [Micikevicius+ 2017]
910
911     Pytorch has an automatic mixed precision (AMP) library.
912     https://pytorch.org/docs/stable/amp.html
913     https://docs.nvidia.com/deeplearning/performance/mixed-precision-training/
914
915     NVIDIA's Transformer Engine supports FP8 for linear layers
916     Use FP8 pervasively throughout training [Peng+ 2023]
917
918
919 #####
920
921 def get_memory_usage(x: torch.Tensor):
922     return x.numel() * x.element_size()
923
924
925 def get_promised_flop_per_sec(device: str, dtype: torch.dtype) -> float:
926     """Return the peak FLOP/s for `device` operating on `dtype`."""
927     if not torch.cuda.is_available():
928         No CUDA device available, so can't get FLOP/s.
929         return 1
930     properties = torch.cuda.get_device_properties(device)
931
932     if "A100" in properties.name:
933         # https://www.nvidia.com/content/dam/en-zz/Solutions/Data-Center/a100/pdf/nvidia-a100-datasheet-us-nvidia-1758950-r4-web.pdf)
934         if dtype == torch.float32:
935             return 19.5e12
936         if dtype in (torch.bfloat16, torch.float16):
937             return 312e12
938         raise ValueError(f"Unknown dtype: {dtype}")
939
940     if "H100" in properties.name:
941         # https://resources.nvidia.com/en-us-tensor-core/nvidia-tensor-core-gpu-datasheet)
942         if dtype == torch.float32:
943             return 67.5e12
944         if dtype in (torch.bfloat16, torch.float16):
945             return 1979e12 / 2 # 1979 is for sparse, dense is half of that
946         raise ValueError(f"Unknown dtype: {dtype}")
947
948     raise ValueError(f"Unknown device: {device}")
949
950
951 def same_storage(x: torch.Tensor, y: torch.Tensor):
952     return x.untyped_storage().data_ptr() == y.untyped_storage().data_ptr()
953

```



```
954
955 def time_matmul(a: torch.Tensor, b: torch.Tensor) -> float:
956     """Return the number of seconds required to perform `a @ b`."""
957
958     # Wait until previous CUDA threads are done
959     if torch.cuda.is_available():
960         torch.cuda.synchronize()
961
962     def run():
963         # Perform the operation
964         a @ b
965
966         # Wait until CUDA threads are done
967         if torch.cuda.is_available():
968             torch.cuda.synchronize()
969
970     # Time the operation `num_trials` times
971     num_trials = 5
972     total_time = timeit.timeit(run, number=num_trials)
973
974     return total_time / num_trials
975
976
977 def get_num_parameters(model: nn.Module) -> int:
978     return sum(param.numel() for param in model.parameters())
979
980 def get_device(index: int = 0) -> torch.device:
981     """Try to use the GPU if possible, otherwise, use CPU."""
982     if torch.cuda.is_available():
983         return torch.device(f"cuda:{index}")
984     else:
985         return torch.device("cpu")
986
987 if __name__ == "__main__":
988     main()
```