

lecture_14.py



```

1 from dataclasses import dataclass
2 import math
3 import torch
4 import torch.nn as nn
5 from torch.nn.functional import softmax
6 import numpy as np
7 import kenlm
8 import fasttext
9 import itertools
10 import mmh3
11 from bitarray import bitarray
12 from basic_util import count, repeat
13 from file_util import download_file
14 from execute_util import text, image, link
15 from lecture_util import article_link, named_link
16 from references import dolma
17
18 def main():
19     Last lecture: overview of datasets used for training language models
20     • Live service (GitHub) → dump/crawl (GH Archive) → processed data (The Stack)
21     • Processing: HTML to text, language/quality/toxicity filtering, deduplication
22
23     This lecture: deep dive into the mechanics
24     • Algorithms for filtering (e.g., classifiers)
25     • Applications of filtering (e.g., language, quality, toxicity)
26     • Deduplication (e.g., Bloom filters, MinHash, LSH)
27
28     filtering_algorithms()
29     filtering_applications()
30     deduplication()
31
32

```


Summary

- Algorithmic tools: n-gram models (KenLM), classifiers (fastText), importance resampling (DSIR)
- Applications: language identification, quality filtering, toxicity filtering
- Deduplication: hashing scales to large datasets for fuzzy matching
- Now you have the tools (mechanics), just have to spend time with data (intuitions)

```

39 def filtering_algorithms():
40     Algorithmic building block:
41     • Given some target data T and lots of raw data R, find subset T' of R similar to T.
42

```



```

43
44 Desiderata for filtering algorithm:
45 • Generalize from the target data (want T and T' to be different)
46 • Extremely fast (have to run it on R, which is huge)

```

```

47
48     kenlm_main()          # Train n-gram model
49     fasttext_main()       # Train a classifier
50     dsir_main()           # Train bag of n-grams model, do importance resampling
51     filtering_summary()
52

```

Survey paper on data selection [\[Albalak+ 2024\]](#)

```

56 def kenlm_main():
57     n-gram model with Kneser-Ney smoothing \[article\]
58     • KenLM: fast implementation originally for machine translation \[code\]
59     • Common language model used for data filtering
60     • Extremely simple / fast - just count and normalize
61
62

```

Concepts

Maximum likelihood estimation of n-gram language model:

- $n = 3$: $p(\text{in} \mid \text{the cat}) = \text{count}(\text{the cat in}) / \text{count}(\text{the cat})$

Problem: sparse counts (count of many n-grams is 0 for large n)

Solution: Use Kneser-Ney smoothing to handle unseen n-grams [\[article\]](#)

- $p(\text{in} \mid \text{the cat})$ depends on $p(\text{in} \mid \text{cat})$ too

```

69     # Download a KenLM language model
70     model_url = "https://huggingface.co/edugp/kenlm/resolve/main/wikipedia/en.arpa.bin"
71     model_path = "var/en.arpa.bin"
72     download_file(model_url, model_path)
73     model = kenlm.Model(model_path)
74

```

Use the language model

```

76 def compute(content: str):
77     # Hacky preprocessing
78     content = "<s> " + content.replace(",", " ,").replace(".", " .") + " </s>"
79

```

log p(content)

```

81     score = model.score(content)
82

```

Perplexity normalizes by number of tokens to avoid favoring short documents

```

84     num_tokens = len(list(model.full_scores(content)))
85     perplexity = math.exp(-score / num_tokens)
86

```

```

87     return score, perplexity
88

```

```

89     score, perplexity = compute("Stanford University was founded in 1885 by Leland and Jane Stanford as a tribute to the
memory of their only child, Leland Stanford Jr.") # @inspect score, @inspect perplexity
90     score, perplexity = compute("If you believe that the course staff made an objective error in grading, you may submit
a regrade request on Gradescope within 3 days after the grades are released.") # @inspect score, @inspect perplexity
91     score, perplexity = compute("asdf asdf asdf asdf asdf") # @inspect score, @inspect perplexity
92     score, perplexity = compute("the the the the the the the the the the the the the the the the") # @inspect score,
@inspect perplexity
93
94

```

CCNet

[\[Wenzek+ 2019\]](#)

- Items are paragraphs of text
- Sort paragraphs by increasing perplexity
- Keep the top 1/3
- Was used in LLaMA

Summary: Kneser-Ney n-gram language models (with KenLM implementation) is fast but crude

```

104 def fasttext_main():

```

fastText classifier [Joulin+ 2016]

- Task: text classification (e.g., sentiment classification)
- Goal was to train a fast classifier for text classification
- They found it was as good as much slower neural network classifiers

Baseline: bag of words (not what they did)

```

111 L = 32                                # Length of input
112 V = 8192                             # Vocabulary size
113 K = 64                                # Number of classes
114 W = nn.Embedding(V, K)                # Embedding parameters (V x K)
115 x = torch.randint(V, (L,))            # Input tokens (L) - e.g., ["the", "cat", "in", "the", "hat"]
116 y = softmax(W(x).mean(dim=0))         # Output probabilities (K)
117 Problem: V*K parameters (could be huge)

```

fastText classifier: bag of word embeddings

```

120 H = 16                                # Hidden dimension
121 W = nn.Embedding(V, H)                # Embedding parameters (V x H)
122 U = nn.Linear(H, K)                  # Head parameters (H x K)
123 y = softmax(U(W(x).mean(dim=0)))     # Output probabilities (K)
124 Only H*(V + K) parameters

```

Implementation:

- Parallelized, asynchronous SGD
- Learning rate: linear interpolation from [some number] to 0 [article]

Bag of n-grams

```

131 x = ["the cat", "cat in", "in the", "the hat"] # @inspect x
132 Problem: number of bigrams can get large (and also be unbounded)
133 Solution: hashing trick
134 num_bins = 8 # In practice, 10M bins
135 hashed_x = [hash(bigram) % num_bins for bigram in x] # @inspect hashed_x

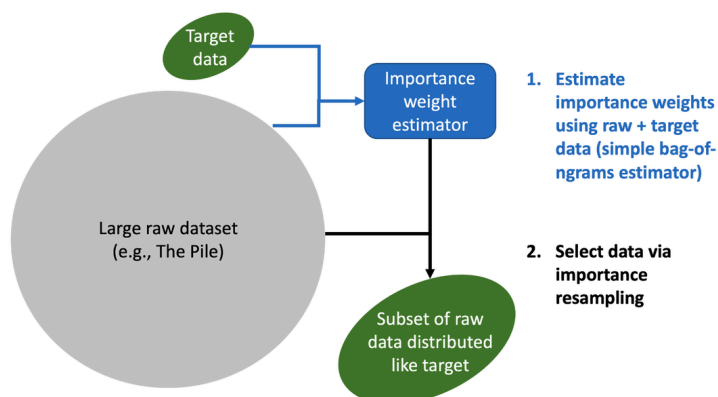
```

- For quality filtering, we have $K = 2$ classes (good versus bad)
- In that case, fastText is just a linear classifier ($H = K = 2$)

In general, can use any classifier (e.g., BERT, Llama), it's just slower

```
143 def dsir_main():
```

Data Selection for Language Models via Importance Resampling (DSIR) [Xie+ 2023]



```
146 importance_sampling()
```

Setup:

- Target dataset D_p (small)
- Proposal (raw) dataset D_q (large)

Take 1:

- Fit target distribution p to D_p
 - Fit proposal distribution q to D_q
 - Do importance resampling with p , q , and raw samples D_q
- Problem: target data D_p is too small to estimate a good model

Take 2: use hashed n-grams

```
training_text = "the cat in the hat"
```

```
# Hash the n-grams
```

```
num_bins = 4
```

```
def get_hashed_ngrams(text: str):
```

```
    ngrams = text.split(" ") # Unigram for now
```

```
    return [hash(ngram) % num_bins for ngram in ngrams]
```

```
training_hashed_ngrams = get_hashed_ngrams(training_text) # @inspect training_hashed_ngrams
```

```
# Learn unigram model
```

```
probs = [count(training_hashed_ngrams, x) / len(training_hashed_ngrams) for x in range(num_bins)] # @inspect probs
```

```
# Evaluate probability of any sentence
```

```
hashed_ngrams = get_hashed_ngrams("the text") # @inspect hashed_ngrams
```

```
prob = np.prod([probs[x] for x in hashed_ngrams]) # @inspect prob
```

Result: DSIR slightly better than heuristic classification (fastText) on the [GLUE](#) benchmark

	MNLI	QNLI	QQP	RTE	SST-2	MRPC	CoLA	STS-B	Avg
Random selection	82.63 _{0.41}	86.90 _{0.28}	89.57 _{0.30}	67.37 _{1.69}	90.05 _{0.41}	87.40 _{1.08}	49.41 _{3.67}	88.63 _{0.11}	80.25
Heuristic classification	82.69 _{0.17}	85.95 _{0.79}	89.77 _{0.32}	68.59 _{1.75}	88.94 _{0.98}	86.03 _{0.93}	48.17 _{3.19}	88.62 _{0.22}	79.85
Top- k Heuristic classification	83.34 _{0.22}	88.62 _{0.24}	89.89 _{0.19}	70.04 _{0.99}	91.15 _{0.76}	86.37 _{1.00}	53.02 _{3.56}	89.30 _{0.11}	81.47
DSIR	83.07 _{0.29}	89.11 _{0.14}	89.80 _{0.37}	75.09 _{2.76}	90.48 _{0.57}	87.70 _{0.68}	54.00 _{1.34}	89.17 _{0.13}	82.30
Top- k DSIR	83.39 _{0.06}	88.63 _{0.38}	89.94 _{0.17}	72.49 _{1.29}	91.01 _{0.79}	86.18 _{1.12}	49.90 _{1.10}	89.52 _{0.21}	81.38

Comparison with fastText:

- Modeling distributions is a more principled approach capturing diversity
- Similar computation complexity
- Both can be improved by better modeling

```
def importance_sampling():
```

```
    Setup:
```

- Target distribution p (want samples from here)
- Proposal distribution q (have samples from here)

```
vocabulary = [0, 1, 2, 3]
```

```
p = [0.1, 0.2, 0.3, 0.4]
```

```
q = [0.4, 0.3, 0.2, 0.1]
```

```
# 1. Sample from q
```

```
n = 100
```

```
samples = np.random.choice(vocabulary, p=q, size = n) # @inspect samples
```

```
Samples (q): [0 2 0 2 1 1 0 1 0 1 0 0 0 0 1 1 0 3 1 1 1 1 1 3 1 0 2 0 1 3 0 2 1 1 2 0 0 0 3 2 1 1 0 1 1 0 3 2 0 2 1 0
1 2 2 2 0 0 0 0 2 1 0 2 0 1 3 0 0 0 0 0 0 0 1 1 2 2 2 1 0 1 1 0 1 0 1 0 1 1 3 3 0 1 0 0 2 0]
```

```
# 2. Compute weights over samples (w \propto p/q)
```

```
w = [p[x] / q[x] for x in samples] # @inspect w
```

```
z = sum(w) # @inspect z
```

```
w = [w_i / z for w_i in w] # @inspect w
```

```
# 3. Resample
```

```
samples = np.random.choice(samples, p=w, size=n) # @inspect samples
```

```
Resampled (p): [2 2 1 3 3 2 0 3 3 2 1 0 3 2 0 3 3 1 3 1 3 2 3 2 3 2 3 1 0 3 2 2 2 0 2 1 2 0 3 1 1 1 3 1 3 3 3 1 0 2 3
1 2 1 2 2 2 2 1 1 0 2 1 1 0 2 3 2 1 3 2 3 1 2 3 2 3 2 1 1 3 3 3 1 1 3 2 1 3 1 3 1 0 3 1 1 2 3 2 1]
```

```

209 def filtering_summary():
210     Implementations: KenLM, fastText, DSIR
211
212

```

General framework

```

213     Given target T and raw R, find subset of R similar to T
214     1. Estimate some model based on R and T and derive a scoring function
215     2. Keep examples in R based on their score
216
217

```

Instantiations of the framework

```

218
219     Generative model of T (KenLM):
220     1. score(x) = p_T(x)
221     2. Keep examples x with score(x) >= threshold (stochastically)
222
223     Discriminative classifier (fastText):
224     1. score(x) = p(T | x)
225     2. Keep examples x with score(x) >= threshold (stochastically)
226
227     Importance resampling (DSIR):
228     1. score(x) = p_T(x) / p_R(x)
229     2. Resample examples x with probability proportional to score(x)
230
231

```

```

232 def filtering_applications():
233     The same data filtering machinery can be used for different filtering tasks.
234     language_identification()
235     quality_filtering()
236     toxicity_filtering()
237
238

```

```

239 def language_identification():
240     Language identification: find text of a specific language (e.g., English)
241
242     Why not just go multilingual?
243     • Data: difficult to do curation / processing of high-quality data in any given language
244     • Compute: in compute-limited regime, less compute/tokens dedicated to any given language
245     Models differ on multilinguality:
246     • English was only 30% of BLOOM (was undertrained), English performance suffered [Laurençon+ 2023]
247     • Most frontier models (GPT-4, Claude, Gemini, Llama, Qwen) are heavily multilingual (sufficiently trained)
248
249     fastText language identification [article]
250     • Off-the-shelf classifier
251     • Supports 176 languages
252     • Trained on multilingual sites: Wikipedia, Tatoeba (translation site) and SETimes (Southeast European news)
253

```

```

254     Example: Dolma keeps pages with p(English) >= 0.5 [Soldaini+ 2024]
255

```

```

256     # Download the model
257     model_url = "https://dl.fbaipublicfiles.com/fasttext/supervised-models/lid.176.bin"
258     model_path = "var/lid.176.bin"
259     download_file(model_url, model_path)
260     model = fasttext.load_model(model_path)
261
262     # Make predictions
263     predictions = model.predict(["The quick brown fox jumps over the lazy dog."]) # English @inspect predictions
264     predictions = model.predict(["The quick brown fox jumps over the lazy dog. The quick brown fox jumps over the lazy
dog."]) # Duplicate @inspect predictions
265     predictions = model.predict(["OMG that movie was 🍌🍌! So dope 🤔👍!"]) # Informal English @inspect predictions
266     predictions = model.predict(["Auf dem Wasser zu singen"]) # German @inspect predictions
267     predictions = model.predict(["The quadratic formula is  $x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$ "]) # Latex @inspect predictions
268     predictions = model.predict(["for (int i = 0; i < 10; i++)"]) # C++ @inspect predictions

```

```

269 predictions = model.predict(["Hello!"]) # English @inspect predictions
270 predictions = model.predict(["Bonjour!"]) # French @inspect predictions
271 predictions = model.predict(["Feliz Navidad / Próspero año y felicidad / I wanna wish you a Merry Christmas"]) #
Spanish + English @inspect predictions
272
273 Caveats:
274 • Difficult for short sequences
275 • Difficult for low-resource languages
276 • Could accidentally filter out dialects of English
277 • Hard for similar languages (Malay and Indonesian)
278 • Ill-defined for code-switching (e.g., Spanish + English)
279
280 OpenMathText [Paster+ 2023]
281 • Goal: curate large corpus of mathematical text from CommonCrawl
282 • Use rules to filter (e.g., contains latex commands)
283 • KenLM trained on ProofPile, keep if perplexity < 15000
284 • Trained fastText classifier to predict mathematical writing, threshold is 0.17 if math, 0.8 if no math
285 Result: produced 14.7B tokens, used to train 1.4B models that do better than models trained on 20x data
286
287
288 def quality_filtering():
289     • Some deliberately do not use model-based filtering (C4, Gopher, RefinedWeb, FineWeb, Dolma)
290     • Some use model-based filtering (GPT-3, LLaMA, DCLM) [becoming the norm]
291
292 GPT-3 [Brown+ 2020]
293 • Positives: samples from {Wikipedia, WebText2, Books1, Books2}
294 • Negatives: samples from CommonCrawl
295
296 Train linear classifier based on word features [article]
297 Keep documents stochastically based on score
298 def keep_document(score: float) -> bool:
299     return np.random.pareto(9) > 1 - score
300
301 ** LLaMA/RedPajama** [Touvron+ 2023]
302 • Positives: samples from pages referenced by Wikipedia
303 • Negatives: samples from CommonCrawl
304 • Keep documents that are classified positive
305
306 phi-1 [Gunasekar+ 2023]
307 Philosophy: really high quality data (textbooks) to train a small model (1.5B)
308 Includes synthetic data from GPT 3.5 (later: GPT-4) and filtered data
309
310 R = "Python subset of the Stack" # Raw data
311 prompt = "determine its educational value for a student whose goal is to learn basic coding concepts"
312 T = "Use GPT-4 with this prompt to classify 100K subset of R to get positive examples"
313 Train random forest classifier on T using output embedding from pretrained codegen model
314 Select data from R that is classified positive by the classifier
315
316 Result on HumanEval:
317 • Train 1.3B LM on Python subset of The Stack (performance: 12.19% after 96K steps)
318 • Train 1.3B LM on new filtered subset (performance: 17.68% after 36K steps) - better!
319
320
321 @dataclass
322 class Example:
323     text: str
324     label: int
325
326
327 def toxicity_filtering():
328     # WARNING: potentially offensive content below
329     Toxicity filtering in Dolma [Soldaini+ 2024]
330
331 Dataset: Jigsaw Toxic Comments dataset (2018) [dataset]

```

```

332 • Project goal: help people have better discussions online [article]
333 • Data: comments on Wikipedia talk page annotated with {toxic, severe_toxic, obscene, threat, insult,
identity_hate}
334
335 Trained 2 fastText classifiers
336 • hate: positive = {unlabeled, obscene}, negative = all else
337 • NSFW: positive = {obscene}, negative = all else
338
339 # Examples from the dataset: (obscene, text)
340 train_examples = [
341     Example(label=0, text="Are you threatening me for disputing neutrality? I know in your country it's quite common
to bully your way through a discussion and push outcomes you want. But this is not Russia."),
342     Example(label=1, text="Stupid peace of shit stop deleting my stuff asshole go die and fall in a hole go to
hell!"),
343 ]
344
345 # Download model
346 model_url = "https://dolma-
artifacts.org/fasttext_models/jigsaw_fasttext_bigrams_20230515/jigsaw_fasttext_bigrams_nsfw_final.bin"
347 model_path = "var/jigsaw_fasttext_bigrams_nsfw_final.bin"
348 download_file(model_url, model_path)
349 model = fasttext.load_model(model_path)
350
351 # Make predictions
352 predictions = model.predict([train_examples[0].text]) # @inspect predictions
353 predictions = model.predict([train_examples[1].text]) # @inspect predictions
354 predictions = model.predict(["I love strawberries"]) # @inspect predictions
355 predictions = model.predict(["I hate strawberries"]) # @inspect predictions
356
357
358 def print_predict(model, content):
359     """Run classifier `model` on `content` and print out the results."""
360     predictions = model.predict([content])
361     print(predictions)
362     #labels, prob =
363     #labels = ", ".join(labels)
364     #text(f"{content} => {labels} {prob}")
365
366
367 def deduplication():
368     Two types of duplicates:
369     • Exact duplicates (mirror sites, GitHub forks) [Gutenberg mirrors]
370     • Near duplicates: same text differing by a few tokens
371
372     Examples of near duplicates:
373     • Terms of service and licenses [MIT license]
374     • Formulaic writing (copy/pasted or generated from a template)

```

Dataset	Example	Near-Duplicate Example
Wiki-40B	<code>\n_START_ARTICLE\nHum Award for Most Impactful Character \n_START_SECTION\nWinners and nominees\n_START_PARAGRAPH\nIn the list below, winners are listed first in the colored row, followed by the other nominees. [...]</code>	<code>\n_START_ARTICLE\nHum Award for Best Actor in a Negative Role \n_START_SECTION\nWinners and nominees\n_START_PARAGRAPH\nIn the list below, winners are listed first in the colored row, followed by the other nominees. [...]</code>
LM1B	<code>I left for California in 1979 and tracked Cleveland's changes on trips back to visit my sisters .</code>	<code>I left for California in 1979 , and tracked Cleveland's changes on trips back to visit my sisters .</code>
C4	<code>Affordable and convenient holiday flights take off from your departure country, "Canada". From May 2019 to October 2019, Condor flights to your dream destination will be roughly 6 a week! Book your Halifax (YHZ) - Basel (BSL) flight now, and look forward to your "Switzerland" destination!</code>	<code>Affordable and convenient holiday flights take off from your departure country, "USA". From April 2019 to October 2019, Condor flights to your dream destination will be roughly 7 a week! Book your Maui Kahului (OGG) - Dubrovnik (DBV) flight now, and look forward to your "Croatia" destination!</code>

```

375 • Minor formatting differences in copy/pasting
376
377 Product description repeated 61,036 times in C4
378 "by combining fantastic ideas, interesting arrangements, and follow the current trends in the field of that
make you more inspired and give artistic touches. We'd be honored if you can apply some or all of these
design in your wedding. believe me, brilliant ideas would be perfect if it can be applied in real and make the
people around you amazed!
379 [example page]
380

```

Deduplication training data makes language models better [\[Lee+ 2021\]](#)

- Train more efficiently (because have fewer tokens)
- Avoid memorization (can mitigate copyright, privacy concerns)

Design space:

1. What is an item (sentence, paragraph, document)?
2. How to match (exact match, existence of common subitem, fraction of common subitems)?
3. What action to take (remove all, remove all but one)?

Key challenge:

- Deduplication is fundamentally about comparing items to other items
- Need linear time algorithms to scale

hash_functions()

exact_deduplication()

bloom_filter()

jaccard_minhash()

locality_sensitive_hashing()

```
def hash_functions():
```

- Hash function h maps item to a hash value (integer or string)
- Hash value much smaller than item
- Hash collision: $h(x) = h(y)$ for $x \neq y$

Tradeoff between efficiency and collision resistance [\[article\]](#)

- Cryptographic hash functions (SHA-256): collision resistant, slow (used in bitcoin)
- DJB2, MurmurHash, CityHash: not collision resistant, fast (used for hash tables)

We will use MurmurHash:

```
h = mmh3.hash("hello") # @inspect h
```

```
def exact_deduplication():
```

Simple example

1. Item: string
2. How to match: exact match
3. Action: remove all but one

Original items

```
items = ["Hello!", "hello", "hello there", "hello", "hi", "bye"] # @inspect items
```

Compute hash -> list of items with that hash

```
hash_items = itertools.groupby(sorted(items, key=mmh3.hash), key=mmh3.hash)
```

Keep one item from each group

```
deduped_items = [next(group) for h, group in hash_items] # @inspect deduped_items
```

- Pro: simple, clear semantics, high precision
- Con: does not deduplicate near duplicates
- This code is written in a MapReduce way, can easily parallelize and scale

C4 [\[Raffel+ 2019\]](#)

1. Item: 3-sentence spans
2. How to match: use exact match
3. Action: remove all but one

Warning: when a 3-sentence span is removed from the middle of a document, the resulting document might not be coherent

```
def bloom_filter():
```

Goal: efficient, approximate data structure for testing set membership


```

444
445 Features of Bloom filters
446 • Memory efficient
447 • Can update, but can't delete
448 • If return 'no', definitely 'no'
449 • If return 'yes', most likely 'yes', but small probability of 'no'
450 • Can drive the false positive rate down exponentially with more time/compute
451
452 items = ["the", "cat", "in", "the", "hat"]
453 non_items = ["what", "who", "why", "when", "where", "which", "how"]
454
455 First, make the range of hash function small (small number of bins).
456 m = 8 # Number of bins
457 table = build_table(items, m)
458 for item in items:
459     assert query_table(table, item, m) == 1
460 result = {item: query_table(table, item, m) for item in non_items} # @inspect result
461 num_mistakes = count(result.values(), True) # @inspect num_mistakes
462 false_positive_rate = num_mistakes / (len(items) + num_mistakes) # @inspect false_positive_rate
463 Problem: false positives for small bins
464
465 Naive solution: increase the number of bins
466 Error probability is  $O(1/\text{num\_bins})$ , decreases polynomially with memory
467
468 Better solution: use more hash functions
469 k = 2 # Number of hash functions
470 table = build_table_k(items, m, k)
471 for item in items:
472     assert query_table_k(table, item, m, k) == 1
473 result = {item: query_table_k(table, item, m, k) for item in non_items} # @inspect result
474 num_mistakes = count(result.values(), 1) # @inspect num_mistakes
475 false_positive_rate = num_mistakes / (len(items) + num_mistakes) # @inspect false_positive_rate
476 Reduced the false positive rate!
477
478 false_positive_rate_analysis()
479
480
481 def false_positive_rate_analysis():
482     Assume independence of hash functions and items \[article\]
483     m = 1000 # Number of bins
484     k = 10 # Number of hash functions
485     n = 100 # Number of items we're inserting
486
487     Consider a test input (not in the set) that would hash into a given test bin (say, i).
488     Now consider putting items into the Bloom filter and seeing if it hits i.
489
490     # Insert one item, ask if the test bin  $B(i) = 1$ ?
491     # B: [0 0 1 0 0 0 0 0 0 0] - have to miss 1 time
492     f = 1 / m #  $P[B(i) = 1 \text{ after } 1 \text{ insertion with } 1 \text{ hash function}]$  # @inspect f
493     # B: [0 0 1 0 0 1 0 1 0 0] - have to miss k times
494     f = 1 - (1 - 1 / m) ** k #  $P[B(i) = 1 \text{ after } 1 \text{ insertion with } k \text{ hash functions}]$  # @inspect f
495
496     # Insert n items, ask if the test bin  $B(i) = 1$ ?
497     # Have to miss  $k \cdot n$  times
498     f = 1 - (1 - 1 / m) ** (k * n) #  $P[B(i) = 1 \text{ after } n \text{ insertions for } 1 \text{ hash function}]$  # @inspect f
499     # Get k chances to miss (since test input is hashed k times too)
500     f = f ** k #  $P[B(i) = 1 \text{ after } n \text{ insertions for } k \text{ hash functions}]$  # @inspect f
501
502     Optimal value of k (given fixed m / n ratio) [results in  $f \sim 0.5$ ]
503     k = math.log(2) * m / n # @inspect k
504     Resulting false positive rate (improved)
505     f = 0.5 ** k # @inspect f
506
507     Tradeoff between compute (k), memory (m), and false positive rate (f) \[lecture notes\]

```

```

508
509 Example: Dolma
510 • Set false positive rate to 1e-15
511 • Perform on items = paragraphs
512
513
514 def build_table(items: list[str], num_bins: int):
515     """Build a Bloom filter table of size `num_bins`, inserting `items` into it."""
516     table = bitarray(num_bins) # @inspect table
517     for item in items:
518         h = mmh3.hash(item) % num_bins # @inspect item, @inspect h
519         table[h] = 1 # @inspect table
520     return table
521
522
523 def build_table_k(items: list[str], num_bins: int, k: int):
524     """Build a Bloom filter table of size `num_bins`, inserting `items` into it.
525     Use `k` hash functions."""
526     table = bitarray(num_bins) # @inspect table
527     for item in items:
528         # For each of the k functions
529         for seed in range(k):
530             h = mmh3.hash(item, seed) % num_bins # @inspect item, @inspect h, @inspect seed
531             table[h] = 1 # @inspect table
532     return table
533
534
535 def query_table(table: bitarray, item: str, num_bins: int, seed: int = 0):
536     """Return whether `item` is in the `table`."""
537     h = mmh3.hash(item, seed) % num_bins
538     return table[h]
539
540
541 def query_table_k(table: bitarray, item: str, num_bins: int, k: int):
542     """Return 1 if table set to 1 for all `k` hash functions."""
543     return int(all(
544         query_table(table, item, num_bins, seed)
545         for seed in range(k)
546     ))
547
548
549 def jaccard_minhash():
550     Let's now look at approximate set membership.
551     First we need a similarity measure.
552
553

```

Jaccard similarity

```

554 Definition: Jaccard(A, B) = |A intersect B| / |A union B|
555 A = {"1", "2", "3", "4"}
556 B = {"1", "2", "3", "5"}
557
558 def compute_jaccard(A, B):
559     intersection = len(A & B) # @inspect intersection
560     union = len(A | B) # @inspect union
561     return intersection / union
562 jaccard = compute_jaccard(A, B) # @inspect jaccard
563
564 Definition: two documents are near duplicates if their Jaccard similarity >= threshold
565
566 Algorithmic challenge: find near duplicates in linear time
567
568

```


MinHash

```

569 MinHash: a random hash function h so that  $\Pr[h(A) = h(B)] = \text{Jaccard}(A, B)$ 
570
571 Normally, you want different items to hash to different hashes
572 ...but here, you want collision probability to depend on similarity
573
574 def minhash(S: set[str], seed: int):
575     return min(mmh3.hash(x, seed) for x in S)
576
577 Characteristic matrix representation:
578 item | A | B
579 1    | 1 | 1
580 2    | 1 | 1
581 3    | 1 | 1
582 4    | 1 | 0
583 5    | 0 | 1
584
585 Random hash function induces a permutation over items
586 Look at which item is first in A and which item is first in B.
587 Each item has the same probability as being first (min)
588 • If 1, 2, 3 is first, then first in A = first in B.
589 • If 4, 5 is first, then first in A  $\neq$  first in B.
590
591 # Verify MinHash approximates Jaccard as advertised
592 n = 100 # Generate this many random hash functions
593 matches = [minhash(A, seed) == minhash(B, seed) for seed in range(n)]
594 estimated_jaccard = count(matches, True) / len(matches) # @inspect estimated_jaccard
595 assert abs(estimated_jaccard - jaccard) < 0.01
596
597 Now we can hash our items, but a collision doesn't tell us  $\text{Jaccard}(A, B) > \text{threshold}$ .
598
599
600 def locality_sensitive_hashing():
601     Locality sensitive hashing (LSH) [book chapter]
602
603     Suppose we hash examples just one MinHash function
604      $P[A \text{ and } B \text{ collide}] = \text{Jaccard}(A, B)$ 
605     On average, more similar items will collide, but very stochastic...
606
607     Goal: have A and B collide if  $\text{Jaccard}(A, B) > \text{threshold}$ 
608     We have to somehow sharpen the probabilities...
609
610     Solution: use n hash functions
611     Break up into b bands of r hash functions each ( $n = b * r$ )
612
613     n = 12 # Number of hash functions
614     b = 3  # Number of bands
615     r = 4  # Number of hash functions per band
616
617     Hash functions:
618     h1 h2 h3 h4 | h5 h6 h7 h8 | h9 h10 h11 h12
619
620     Key: A and B collide if for some band, all its hash functions return same value
621     As we will see, the and-or structure of the bands sharpens the threshold
622
623     Given  $\text{Jaccard}(A, B)$ , what is the probability that A and B collide?
624
625     def get_prob_collision(sim, b, r): # @inspect sim, @inspect b, @inspect r
626         prob_match = sim ** r # Probability that a fixed band matches @inspect prob_match
627         prob_collision = 1 - (1 - prob_match) ** b # Probability that some band matches @inspect prob_collision
628         return prob_collision
629
630     Example
631     prob_collision = get_prob_collision(sim=0.8, b=5, r=10) # @inspect prob_collision

```

631



632

633

```
634 sims = [0.7, 0.75, 0.8, 0.85, 0.9, 0.95, 0.98]
```

```
635 probs = {sim: get_prob_collision(sim=sim, b=10, r=10) for sim in sims} # @inspect probs
```

636

637 Increasing r sharpens the threshold and moves the curve to the right (harder to match)


```
638 probs = {sim: get_prob_collision(sim=sim, b=10, r=20) for sim in sims} # @inspect probs
```

639

640 Increasing b moves the curve to the left (easier to match)

```
641 probs = {sim: get_prob_collision(sim=sim, b=20, r=20) for sim in sims} # @inspect probs
```

642



643

644 Example setting [Lee+ 2021]: $n = 9000$, $b = 20$, $r = 450$

645 $b = 20$

646 $r = 450$

647 What is the threshold (where the phase transition happens)?

```
648 threshold = (1 / b) ** (1 / r) # @inspect threshold
```

649 Probability that a fixed band matches:

```
650 prob_match = (1 / b) # @inspect prob_match
```

651 Probability that A and B collide ($\approx 1-1/e$):

```
652 prob_collision = 1 - (1 - 1 / b) ** b # @inspect prob_collision
```

653

654

```
655 if __name__ == "__main__":
```

```
656     main()
```